

최근 네트워크 트렌드와 PacketNgin을 이용한 실습 그리고 전망

오픈프론티어 1기


김성민 Semih.Kim@gmail.com

목차

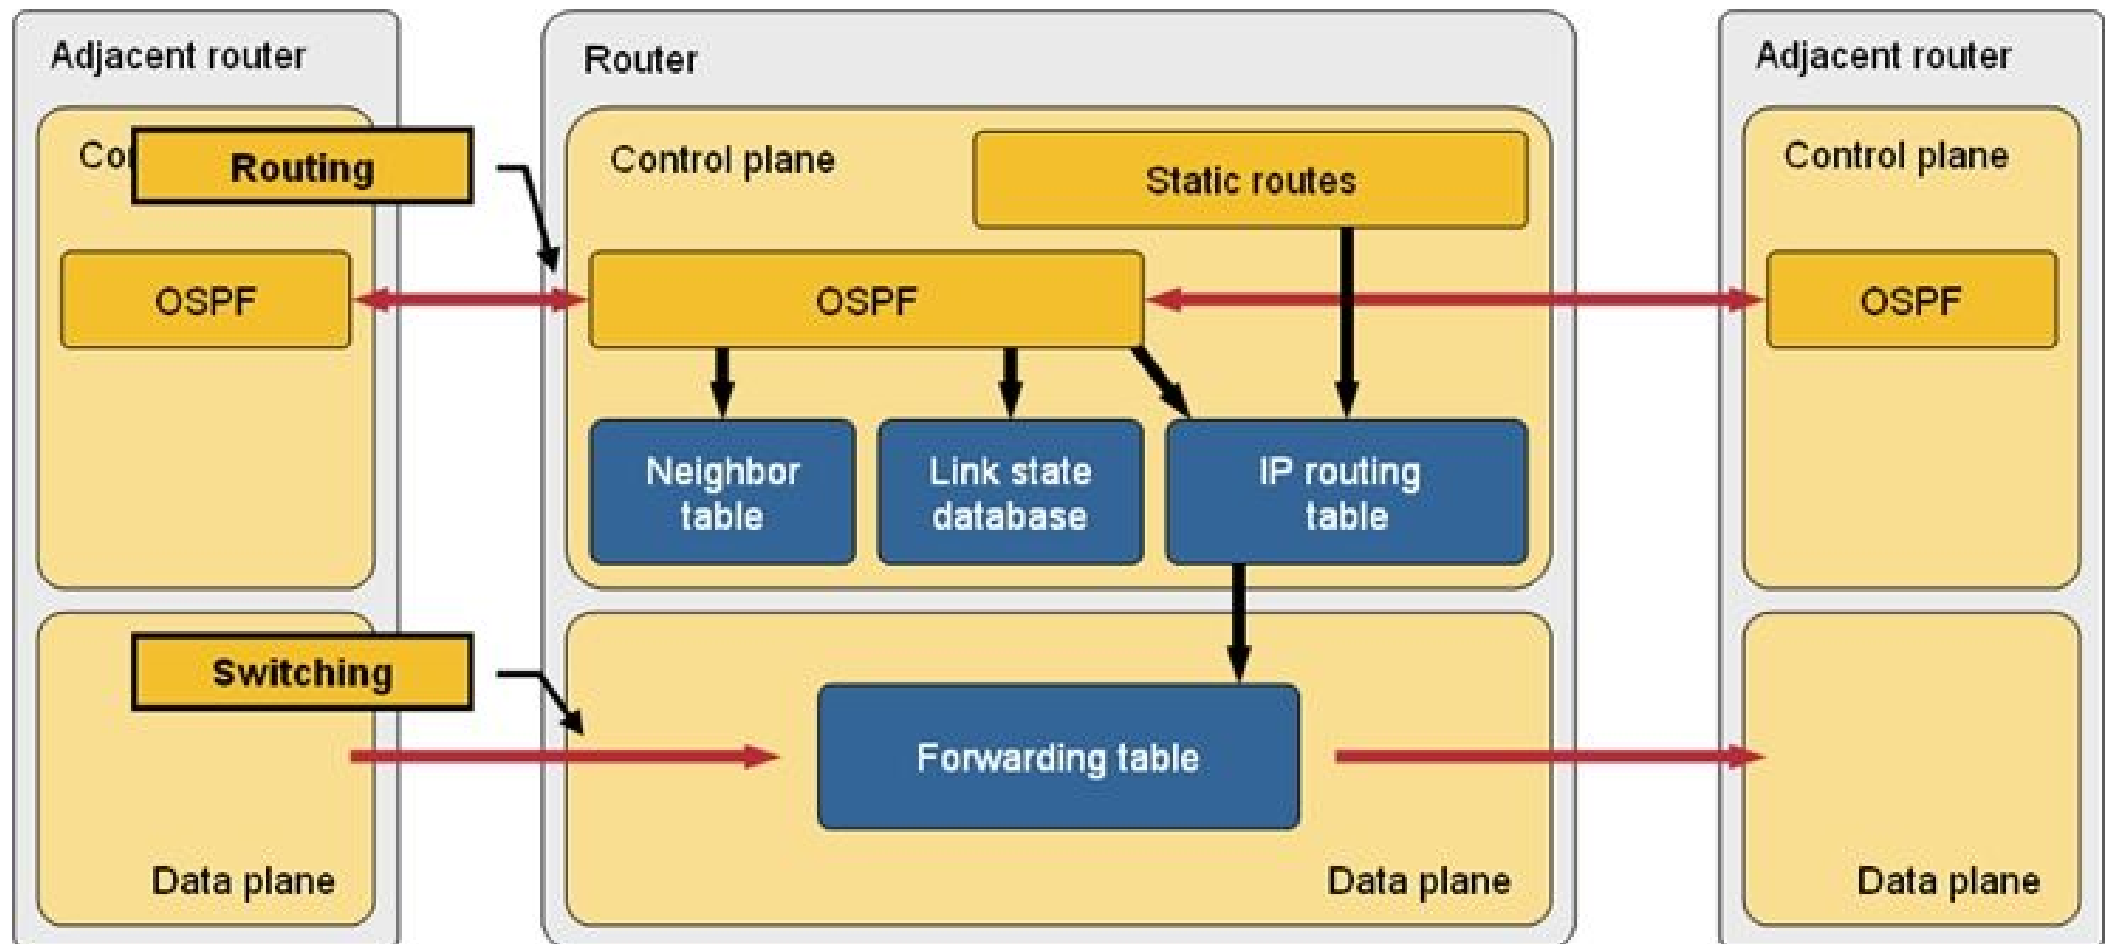
- 트렌드
 - Control and Data plane
 - SDN - Software Defined Network
 - NFV - Network Function Virtualization
- 실습
 - SDN - ONOS
 - NFV - PacketNgin
- 전망
 - OpenFlow의 한계
 - NFV의 한계
 - SDN과 NFV의 결합
 - 말랑말랑한 네트워크

트렌드

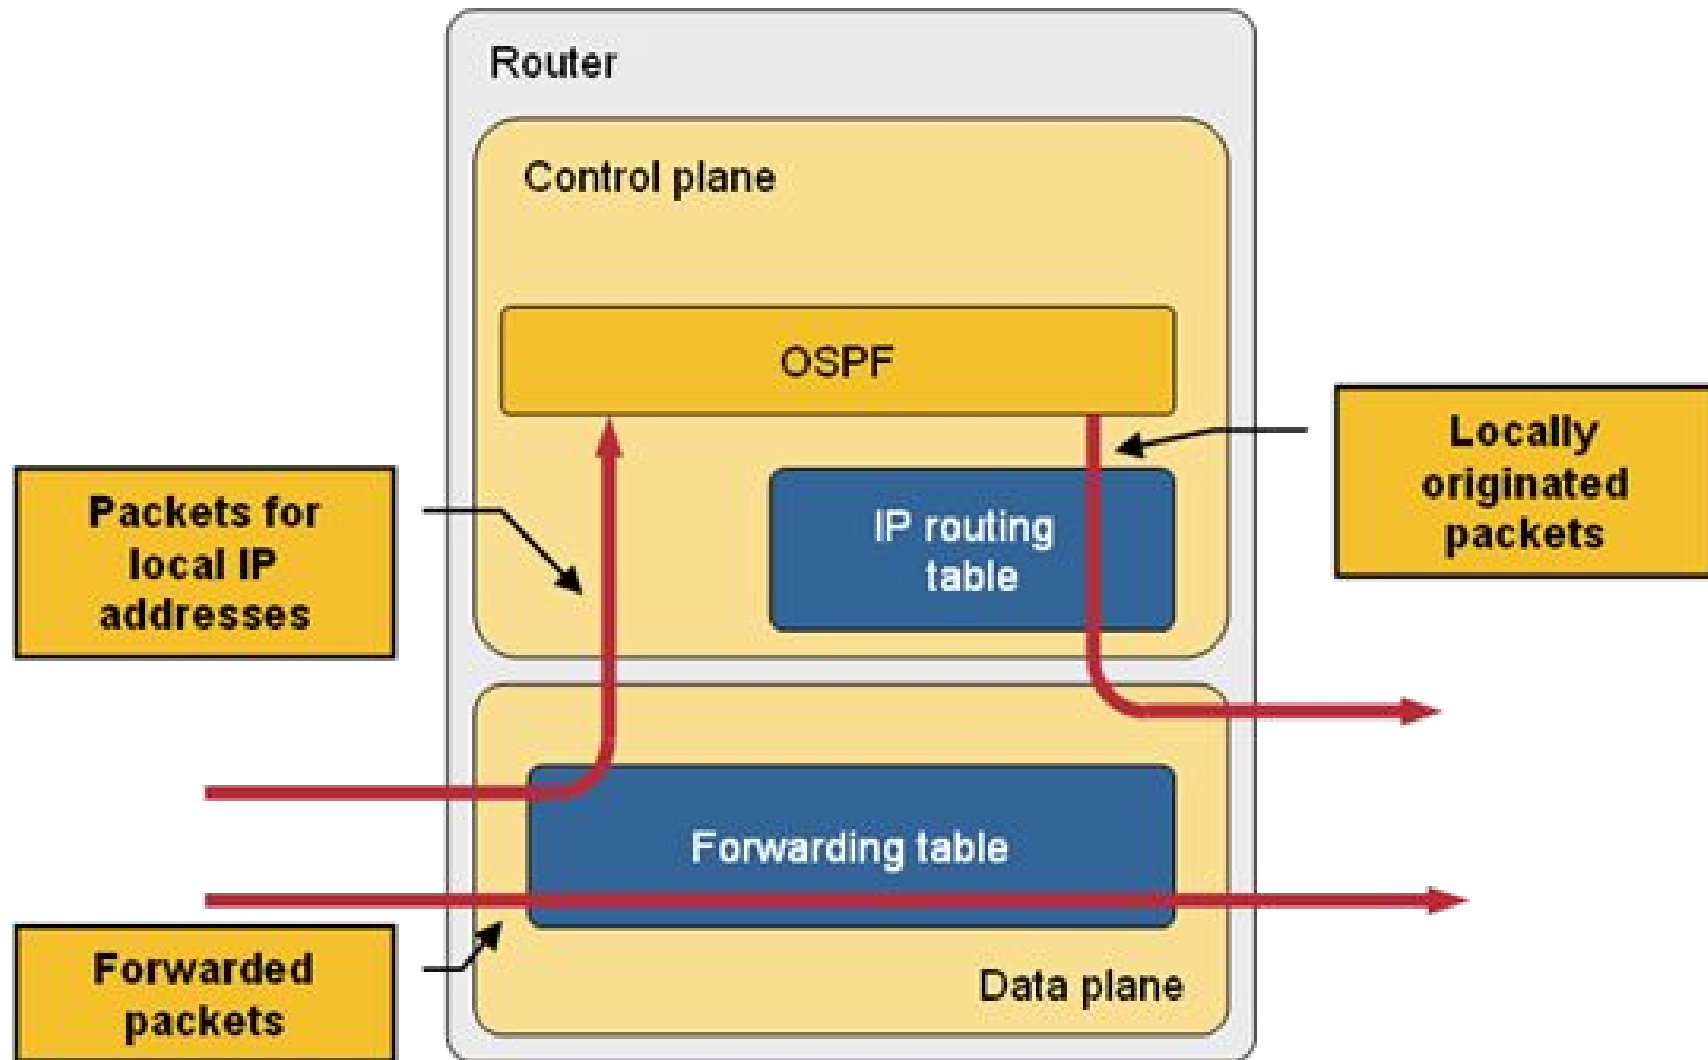
Control and Data plane



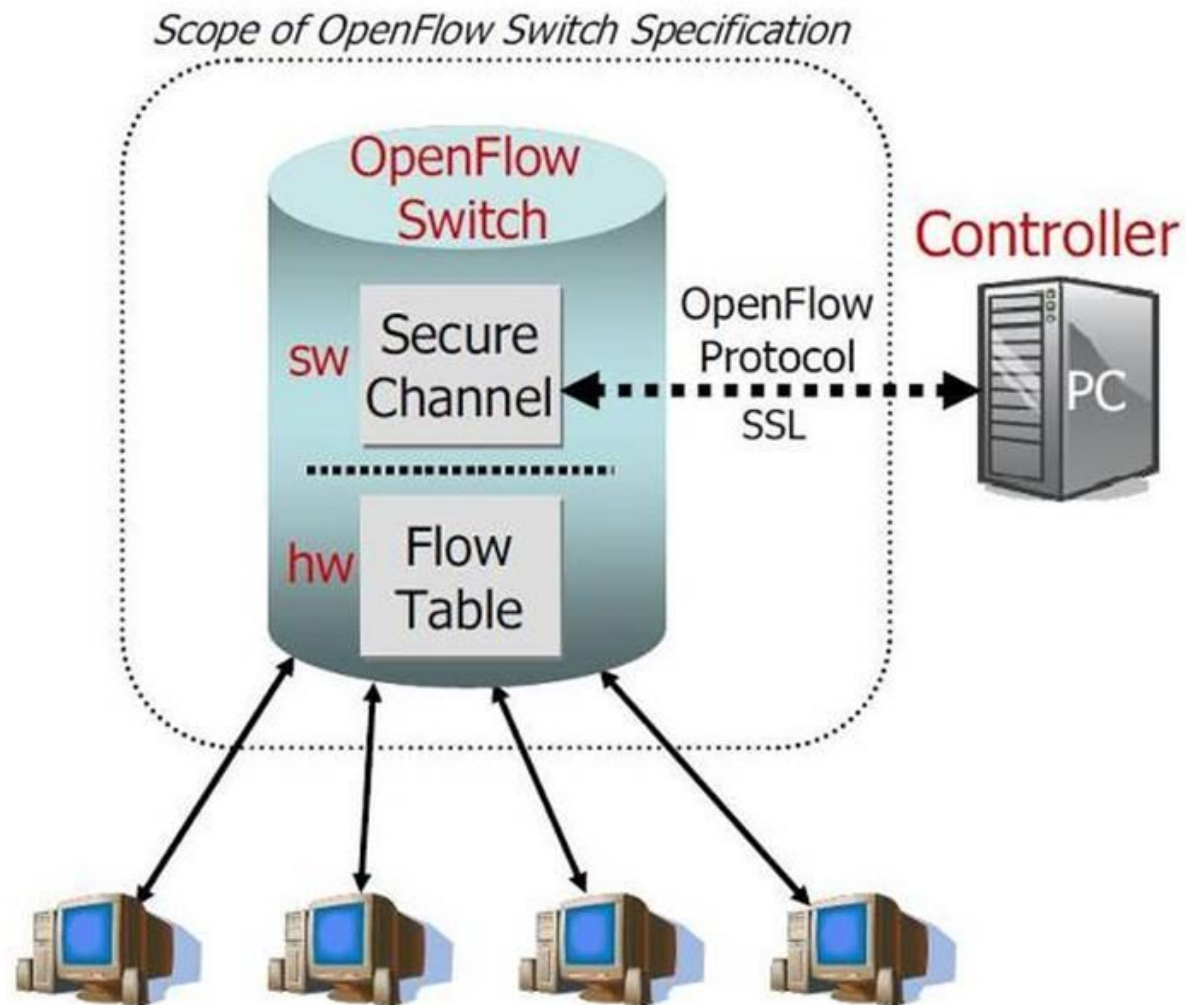
Control and Data plane



Control and Data plane

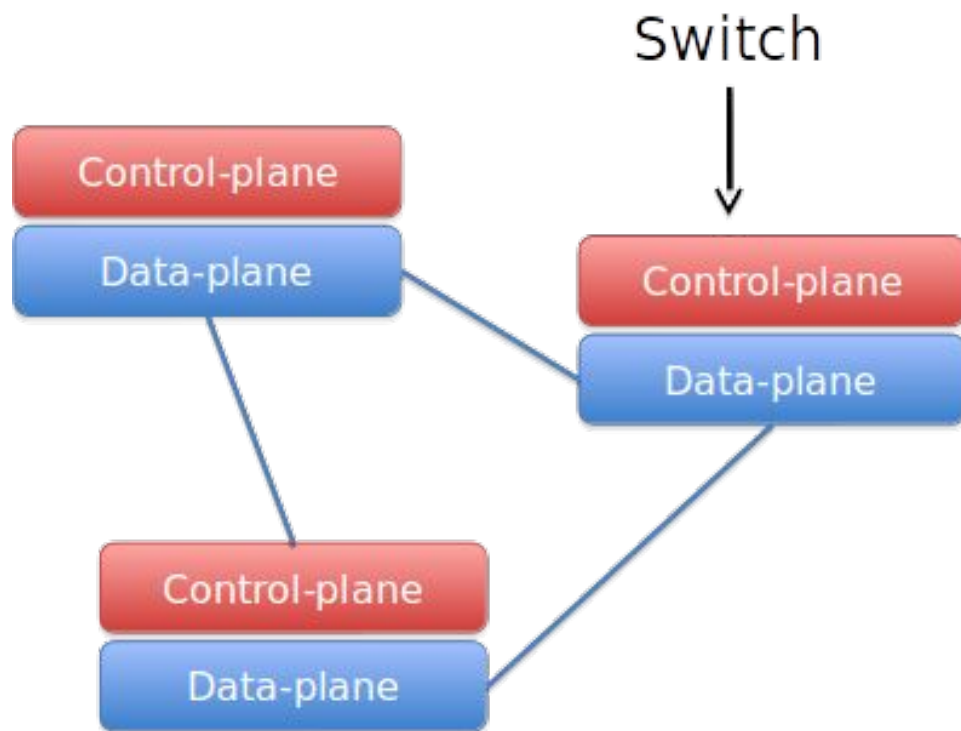


SDN - Software Defined Network




SDN - Software Defined Network

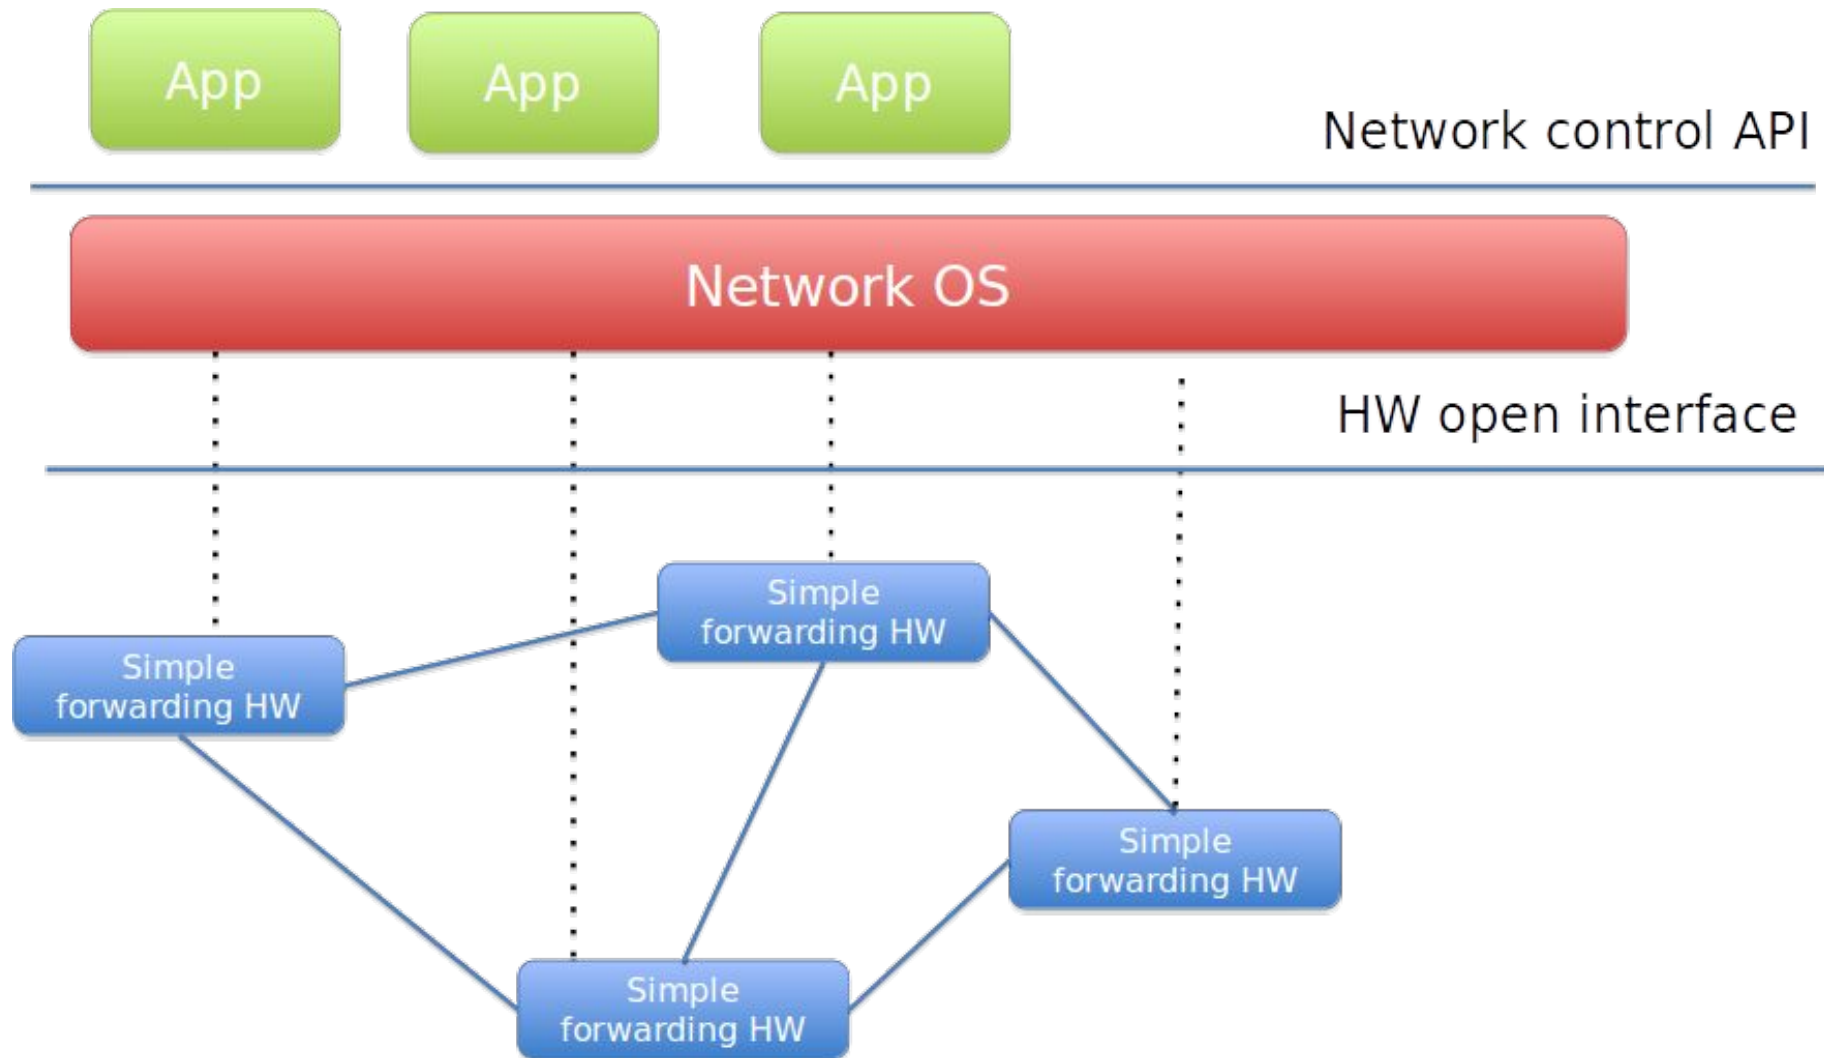
Traditional networking



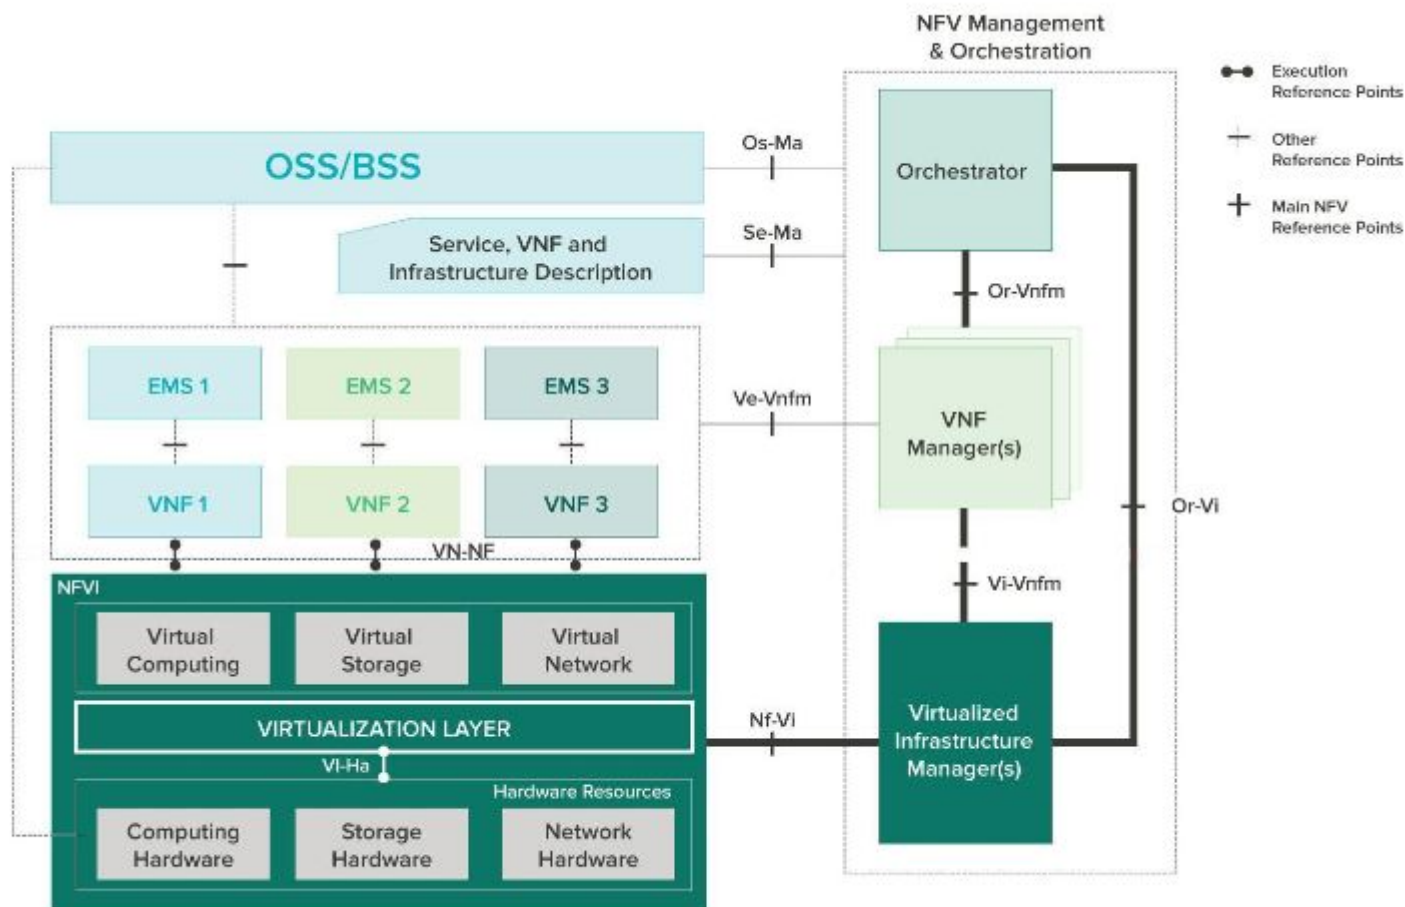
Software-Defined Networking



SDN - Software Defined Network

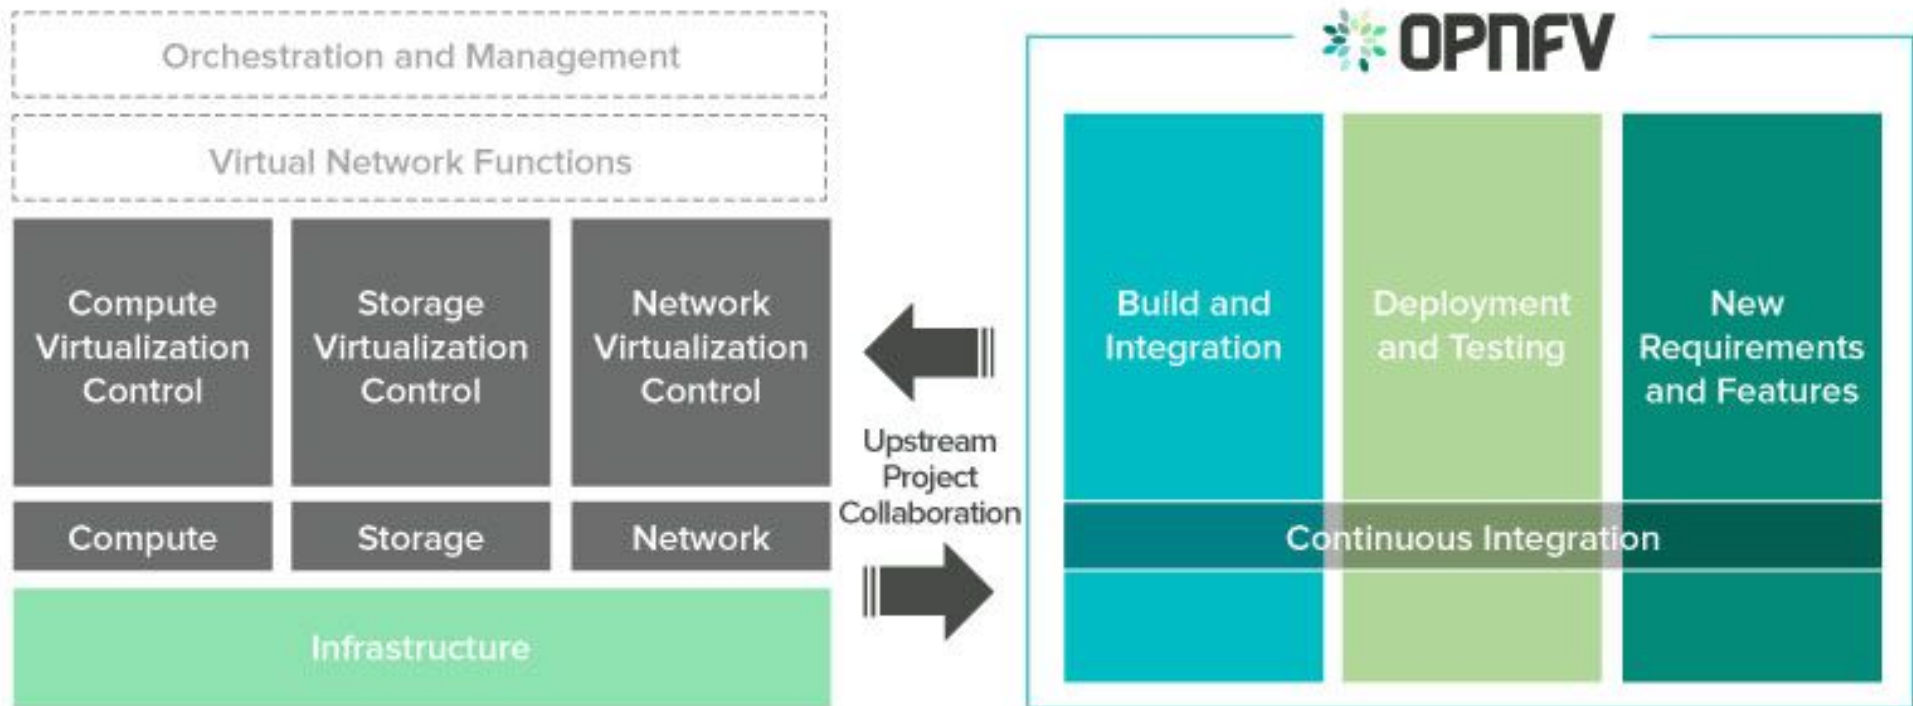


NFV - Network Function Virtualization

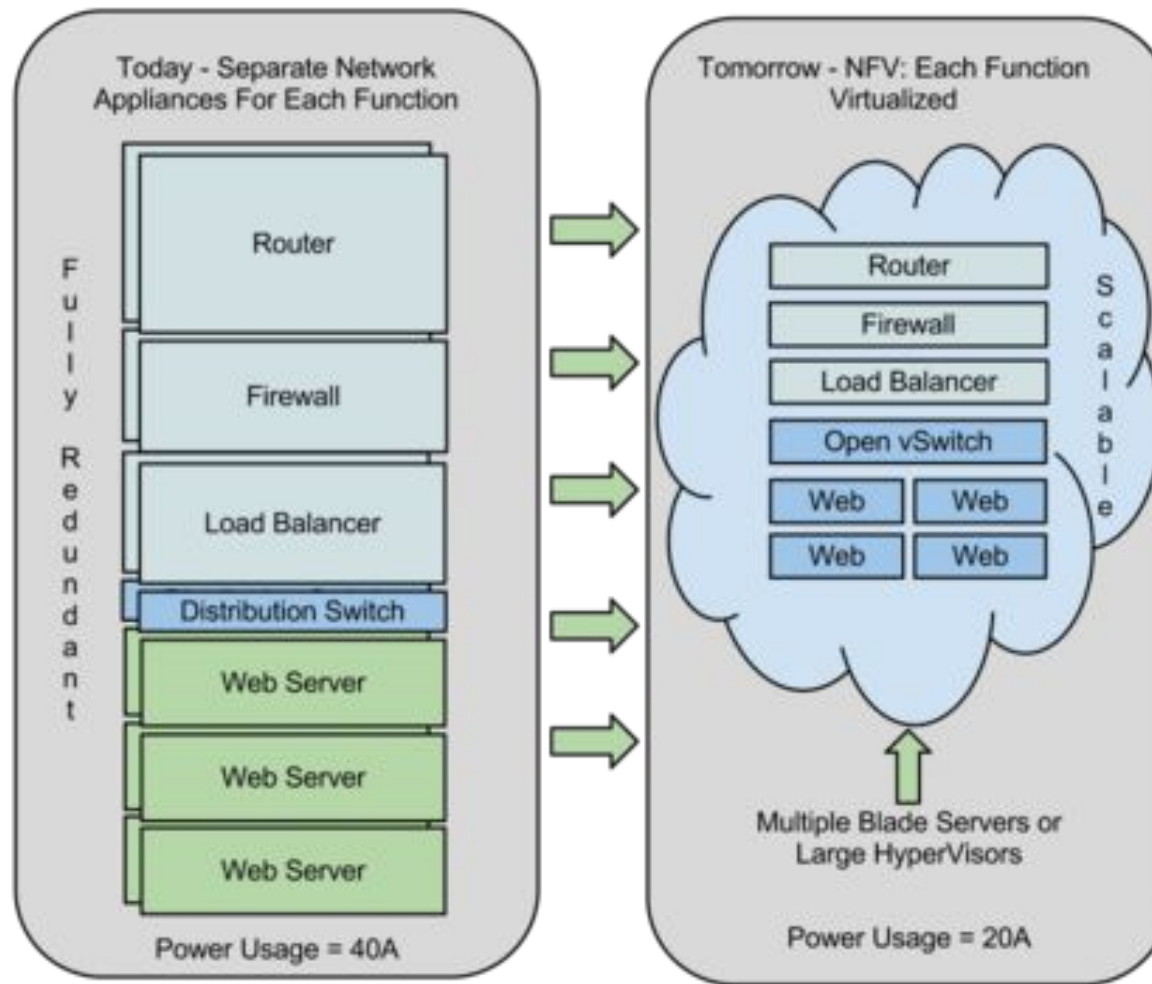


NFV - Network Function Virtualization

OPNFV Platform Overview



NFV - Network Function Virtualization



NFV - Network Function Virtualization

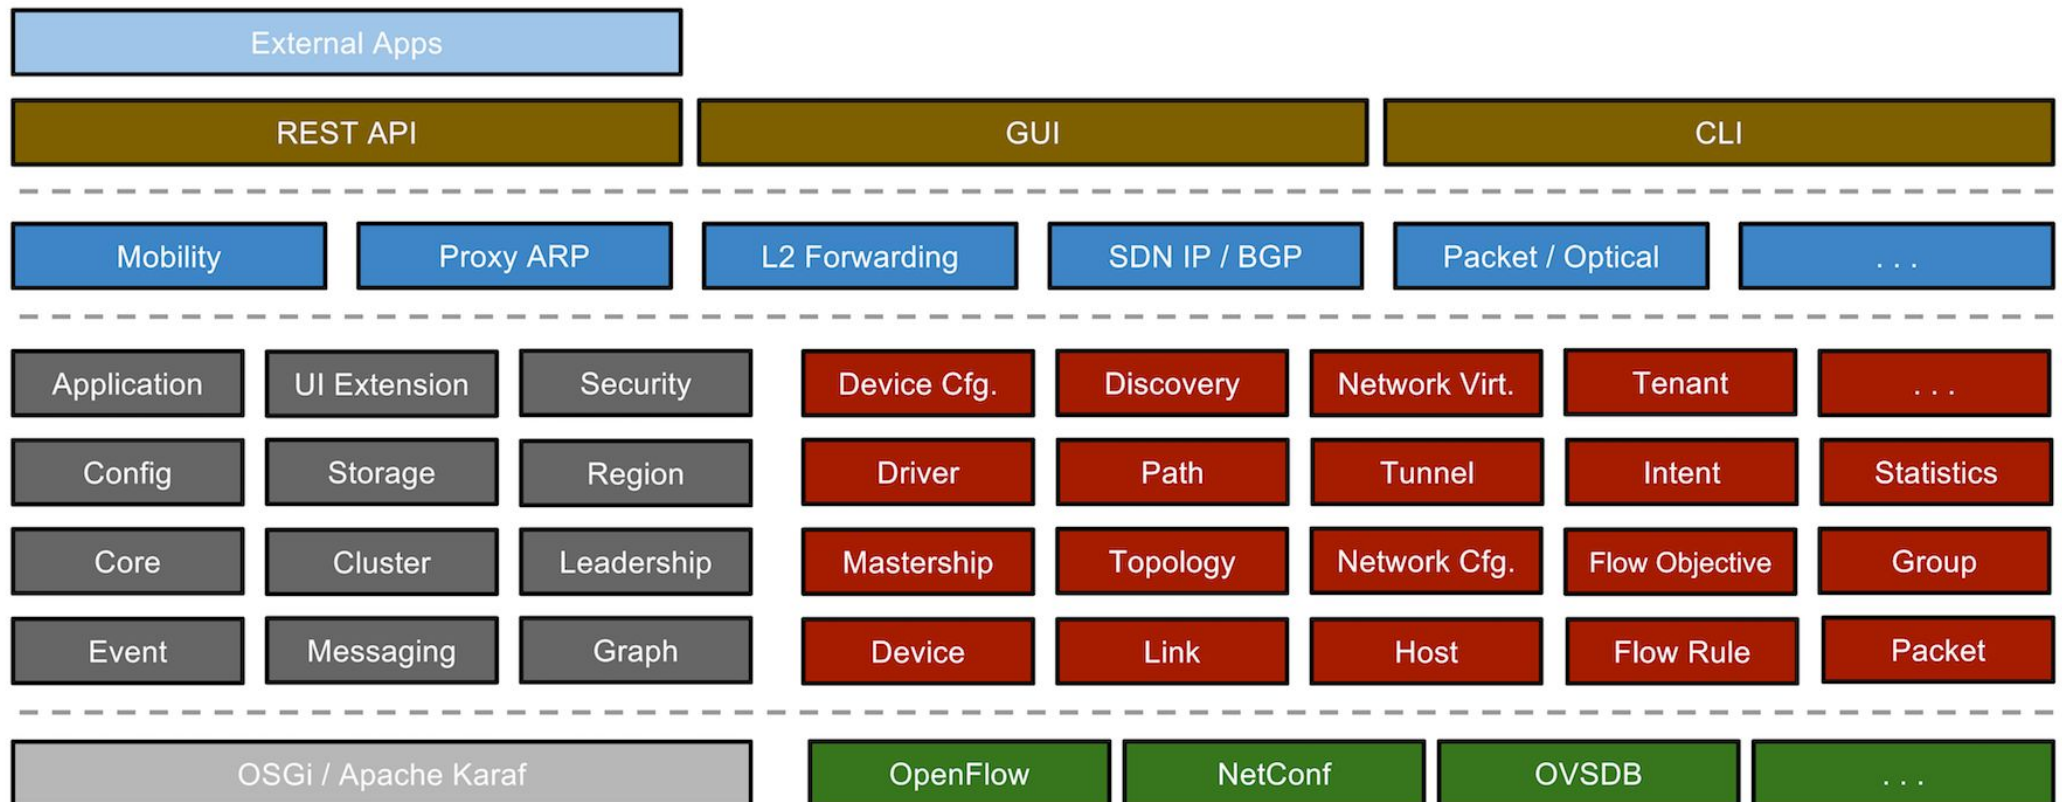
- Mobile Gateway
 - SGW - Serving Gateway
 - PGW - Packet Data Network Gateway
- Tunneling Gateway
 - VPN - Virtual Private Network
- Firewall
 - Web firewall
 - E-mail firewall
 - Intrusion detection system
- Monitoring
 - SLA - Service Level Agreement
- Traffic Control
 - Load balancers
 - WAN accelerators

실습

SDN - ONOS

- Architecture
- Example - ARP Response

ONOS - Architecture



ONOS Example - ARP Response

```
public void programArpRules(DeviceId deviceId, SegmentationId segmentationId,
                           Ip4Address routerAddress, Objective.Operation type) {
    log.info("Install ARP rules");
    TrafficSelector selector = DefaultTrafficSelector.builder()
        .add(Criteria.matchTunnelId(Long
            .parseLong(segmentationId.toString())))
        .matchEthType(EthType.EtherType.ARP.ethType().toShort())
        .matchArpTpa(routerAddress)
        .build();

    TrafficTreatment treatment = DefaultTrafficTreatment.builder().punt()
        .build();

    ForwardingObjective.Builder objective = DefaultForwardingObjective
        .builder().withTreatment(treatment).withSelector(selector)
        .fromApp(appId).withFlag(ForwardingObjective.Flag.SPECIFIC)
        .withPriority(DEFAULT_PRIORITY);

    if (type.equals(Objective.Operation.ADD)) {
        flowObjectiveService.forward(deviceId, objective.add());
    } else {
        flowObjectiveService.forward(deviceId, objective.remove());
    }
}
```

```

private class InternalPacketProcessor implements PacketProcessor {

    @Override
    public void process(PacketContext context) {
        if (context.isHandled()) {
            return;
        }

        InboundPacket pkt = context.inPacket();
        Ethernet ethernet = pkt.parsed();

        if (ethernet == null)
            return;

        if (ethernet.getEtherType() == Ethernet.TYPE_ARP) {

            ARP arp = (ARP) ethernet.getPayload();

            if (arp.getOpCode() == ARP.OP_REQUEST) {
                byte[] srcMacAddress = arp.getSenderHardwareAddress();
                byte[] srcIPAddress = arp.getSenderProtocolAddress();
                byte[] dstIPAddress = arp.getTargetProtocolAddress();

                // Searches the Dst MAC Address based on openstackPortMap
                log.trace("ARP request to : {}", Ip4Address.valueOf(dstIPAddress));

                RouterInterface routerInterface = routerInterfaceStore.values().stream()
                    .filter(e -> {
                        VirtualPort virtualPort = virtualPortService.getPort(e.portId());
                        if (virtualPort == null)
                            return false;

                        for (FixedIp fixedIp : virtualPort.fixedIps()) {
                            IpAddress ipAddress = IpAddress.valueOf(IpAddress.Version.INET, dstIPAddress);
                            if(fixedIp.ip().equals(ipAddress))
                                return true;
                        }

                        return false;
                    })
                    .findAny().orElse(null);
            }
        }
    }
}

```

```

MacAddress macAddress = null;
if (routerInterface != null) {
    VirtualPort virtualPort = virtualPortService.getPort(routerInterface.portId());
    macAddress = virtualPort.macAddress();
    log.info("ARP Repose from : {}", macAddress.toString());
} else {
    context.block();
    return;
}

// Creates a response packet
ARP arpReply = new ARP();
arpReply.setOpCode(ARP.OP_REPLY)
    .setHardwareAddressLength(arp.getHardwareAddressLength())
    .setHardwareType(arp.getHardwareType())
    .setProtocolAddressLength(arp.getProtocolAddressLength())
    .setProtocolType(arp.getProtocolType())
    .setSenderHardwareAddress(macAddress.toBytes())
    .setSenderProtocolAddress(dstIPAddress)
    .setTargetHardwareAddress(srcMacAddress)
    .setTargetProtocolAddress(srcIPAddress);

// Sends a response packet
ethernet.setDestinationMACAddress(srcMacAddress)
    .setSourceMACAddress(macAddress)
    .setEtherType(Ethernet.TYPE_ARP)
    .setPayload(arpReply);


TrafficTreatment.Builder builder = DefaultTrafficTreatment.builder();
builder.setOutput(pkt.receivedFrom().port());
OutboundPacket packet = new DefaultOutboundPacket(pkt.receivedFrom().deviceId(),
    builder.build(), ByteBuffer.wrap(ethernet.serialize()));
packetService.emit(packet);

```

NFV - PacketNgin

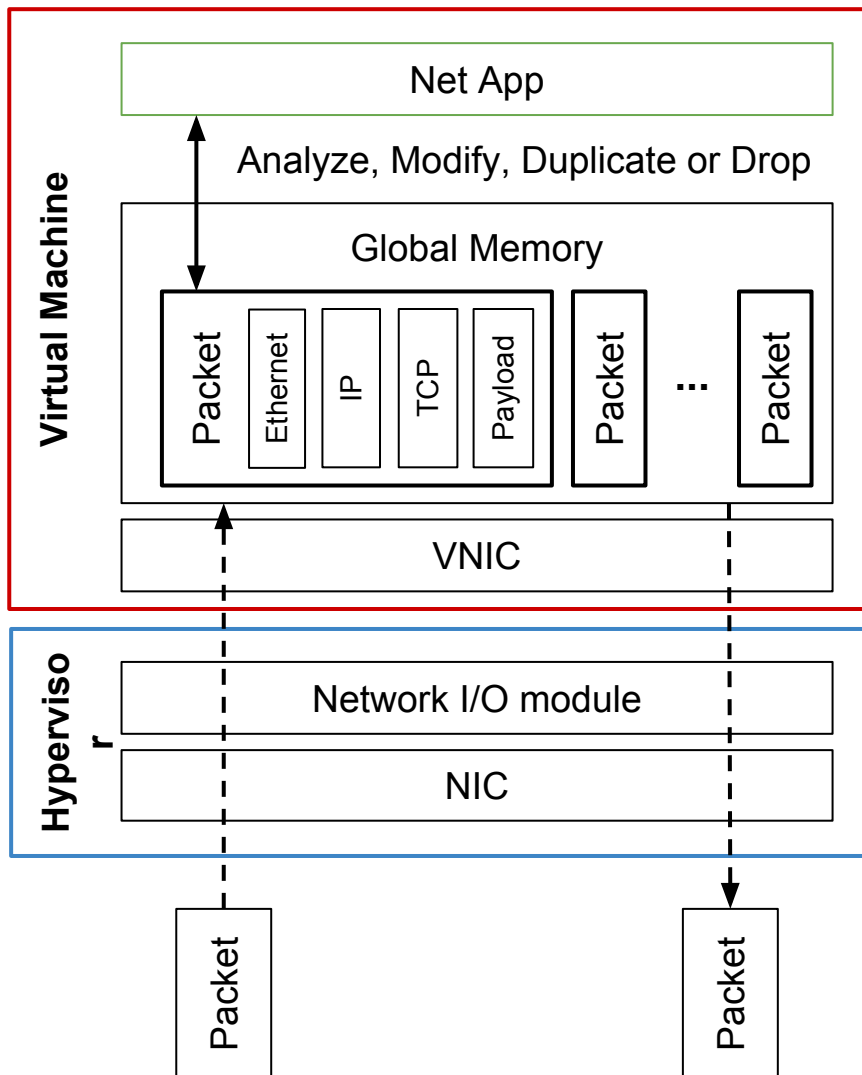
- Concept
- Key Functions
- Example - ARP (Address Resolution Protocol)
- Example - IP (Internet Protocol)
- Example - TCP (Transmission Control Protocol)

PacketNgin Concept



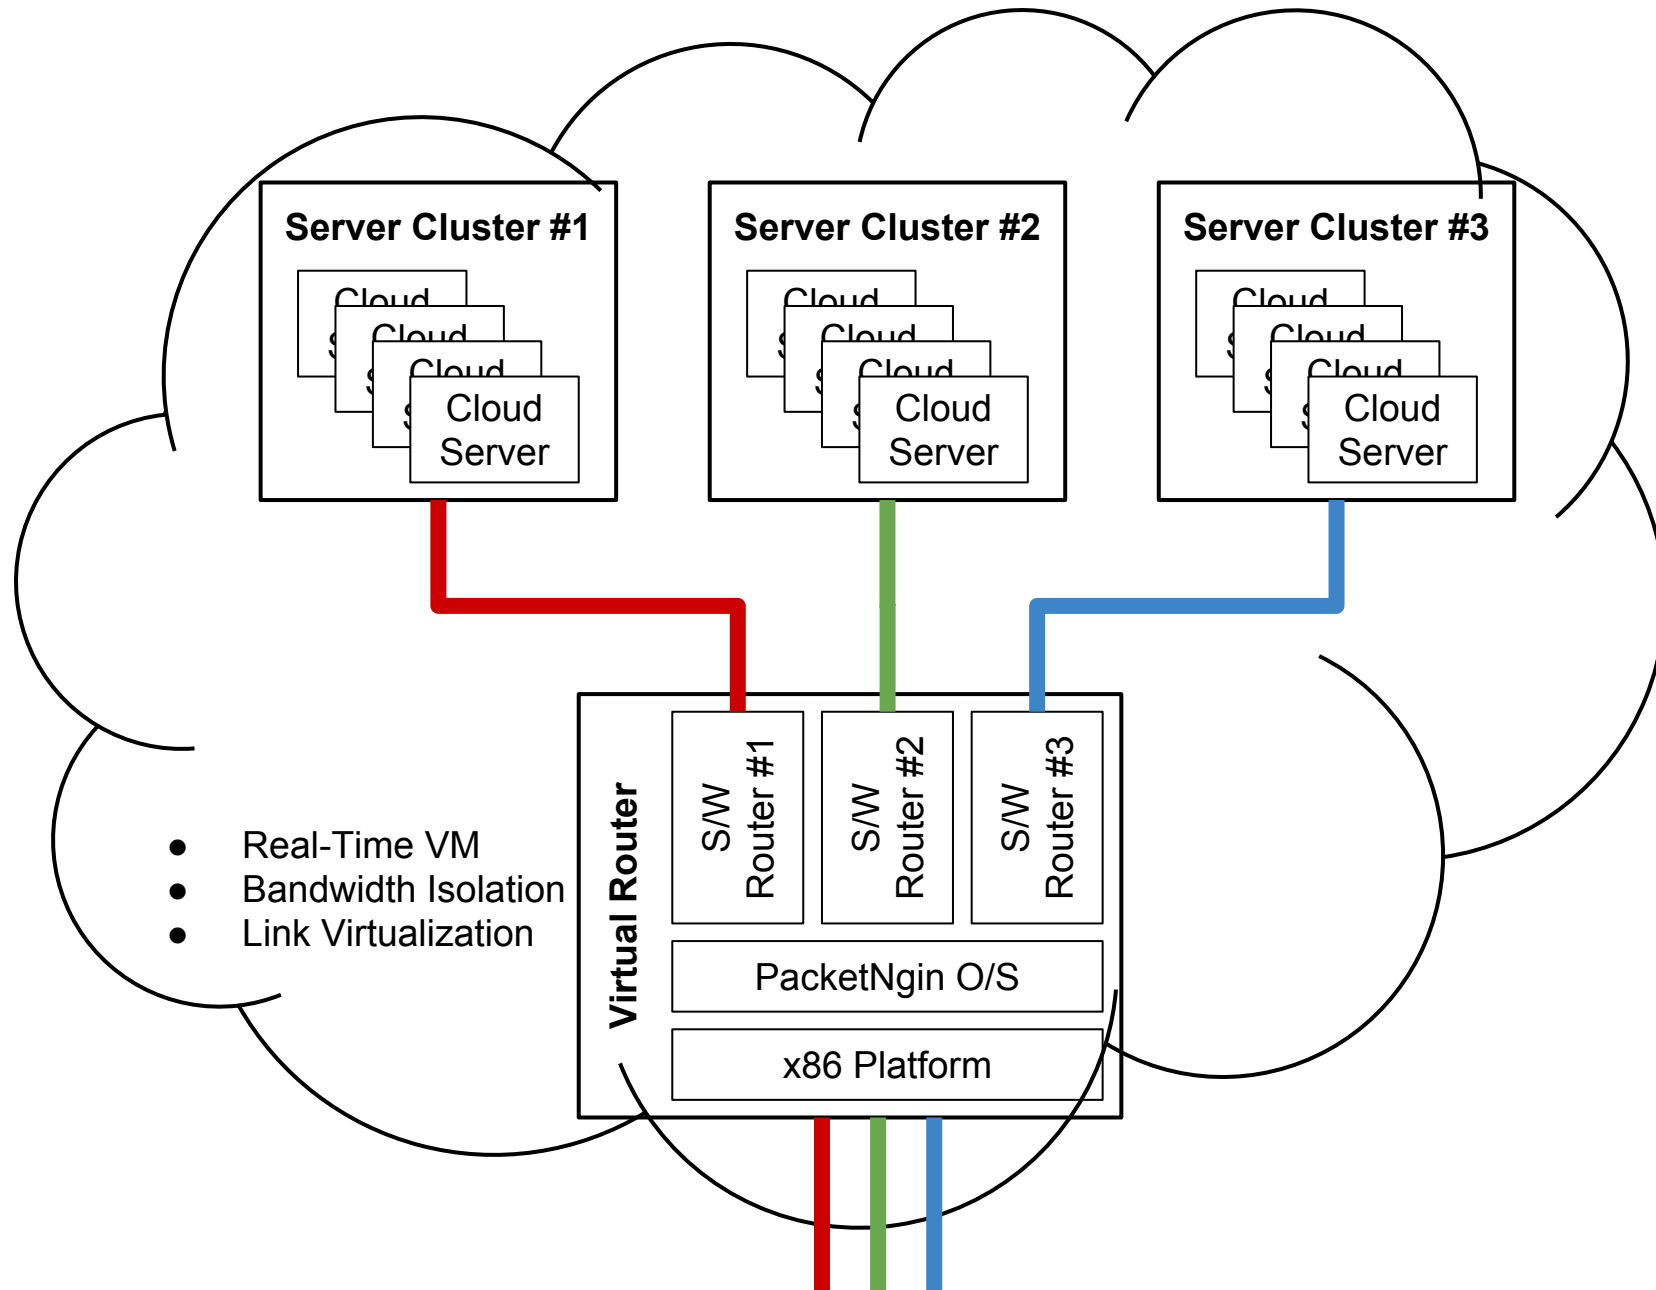
- PacketNgin is an O/S for Network Applications (Net App)
- PacketNgin Virtual Machine is a Real-Time Virtual Machine for Net Apps

PacketNgin - Key Functions

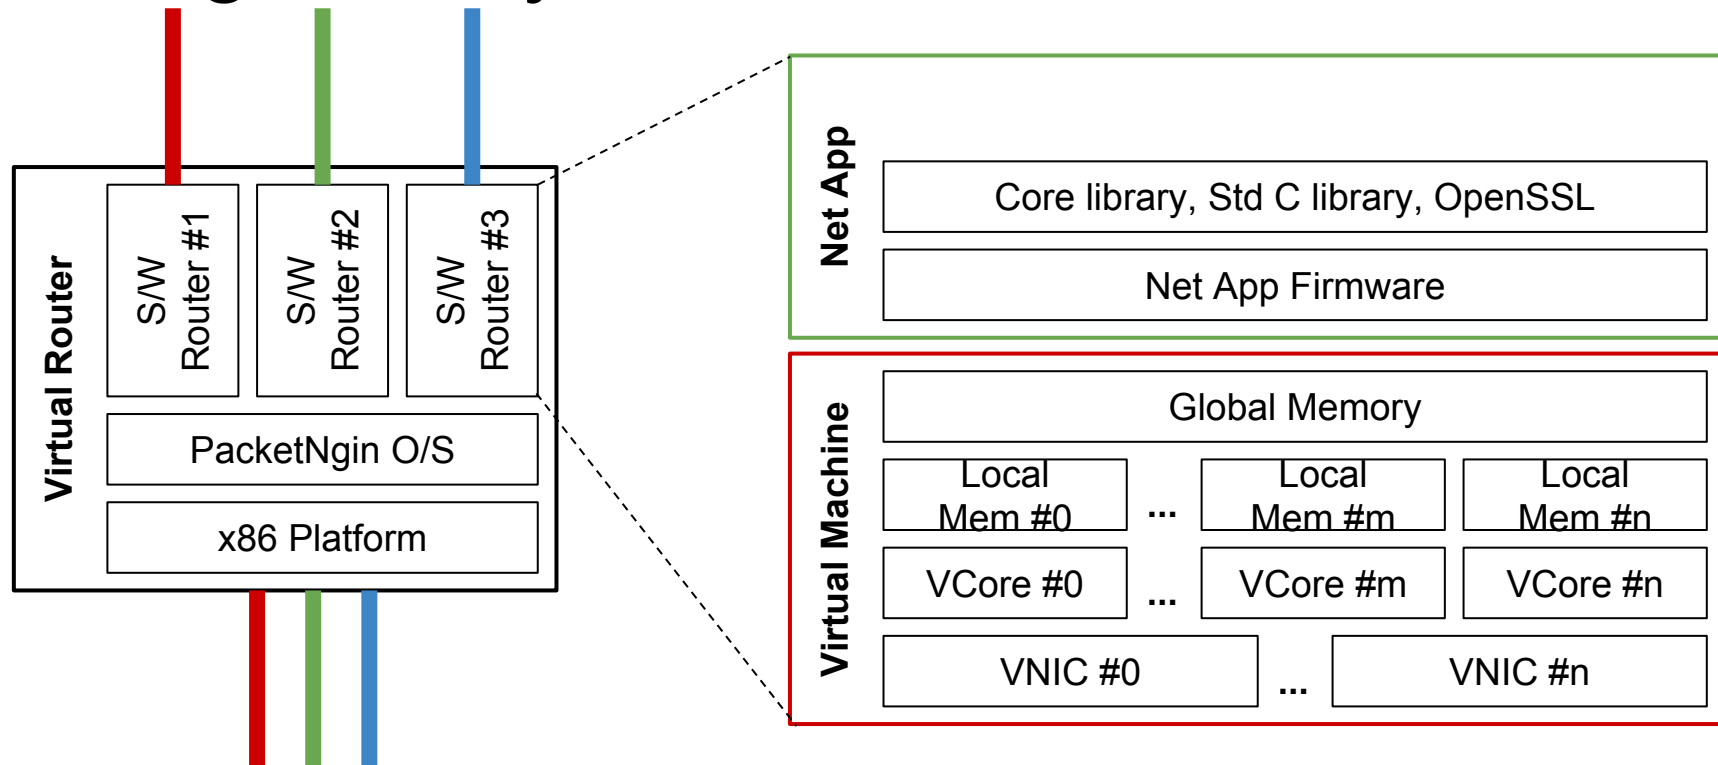


- PacketNgin O/S pass a packet to Net App's memory
- Net App can access the packet from OSI Level 2 to 7 (Ethernet, IP, TCP and payload)
- Net App can analyze, modify, duplicate or drop the packet

PacketNgin - Key Functions




PacketNgin - Key Functions



- PacketNgin O/S creates Real-Time VM
- Real-Time VM consists of
 - dedicated CPU Core, memory area
 - dedicated MAC addr, I/O buffer, bandwidth
 - virtualized link (hop-by-hop or tunnel)

PacketNgin Example - ARP



PacketNgin Example - ARP

```
.657981 ARP, Ethernet (len 6), IPv4 (len 4), Request
0x0000:  ffff ffff ffff bc5f f4bf 0155 0806 0001
0x0010:  0800 0604 0001 bc5f f4bf 0155 cbfc b4b4
0x0020:  0000 0000 0000 cbfc b4fe

.658348 ARP, Ethernet (len 6), IPv4 (len 4), Reply
0x0000:  bc5f f4bf 0155 0000 0c07 acb4 0806 0001
0x0010:  0800 0604 0002 0000 0c07 acb4 cbfc b4fe
0x0020:  bc5f f4bf 0155 cbfc b4b4 0000 0000 0000
0x0030:  0000 0000 0000 0000 0000 0000
```

```
static uint32_t address = 0xc0a8640a;    // 192.168.100.10

void process(NetworkInterface* ni) {
    Packet* packet = ni_input(ni);
    if(!packet)
        return;

    Ether* ether = (Ether*)(packet->buffer + packet->start);

    if(endian16(ether->type) == ETHER_TYPE_ARP) {
        // ARP response
        ARP* arp = (ARP*)ether->payload;
        if(endian16(arp->operation) == 1 && endian32(arp->tpa) == address) {
            ether->dmac = ether->smac;
            ether->smac = endian48(ni->mac);
            arp->operation = endian16(2);
            arp->tha = arp->sha;
            arp->tpa = arp->spa;
            arp->sha = ether->smac;
            arp->spa = endian32(address);


            ni_output(ni, packet);
            packet = NULL;
        }
    }

    if(packet)
        ni_free(packet);
}
```

PacketNgin Example - ARP

```
ubuntu@ubuntu-VirtualBox:~/sdk$ arping 192.168.100.10 -I tap0
ARPING 192.168.100.10 from 192.168.100.200 tap0
Unicast reply from 192.168.100.10 [07:49:CA:5D:F3:EA] 405.598ms
Unicast reply from 192.168.100.10 [07:49:CA:5D:F3:EA] 365.855ms
Unicast reply from 192.168.100.10 [07:49:CA:5D:F3:EA] 282.799ms
Unicast reply from 192.168.100.10 [07:49:CA:5D:F3:EA] 203.837ms
^C Sent 5 probes (1 broadcast(s))
Received 4 response(s)
ubuntu@ubuntu-VirtualBox:~/sdk$
```

PacketNgin Example - IP



Source: http://en.wikipedia.org/wiki/Routing_table

PacketNgin Example - IP

IPv4 Header Format

| Offsets | Octet | 0 | | | | | | | | 1 | | | | | | | | 2 | | | | | | | | 3 | | | | | | | |
|---------|-------|------------------------|---|---|---|-----|---|---|---|----------|---|----|----|----|----|-----|-------|-----------------|-----------------|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| Octet | Bit | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 | 25 | 26 | 27 | 28 | 29 | 30 | 31 |
| 0 | 0 | Version | | | | IHL | | | | DSCP | | | | | | ECN | | Total Length | | | | | | | | | | | | | | | |
| 4 | 32 | Identification | | | | | | | | | | | | | | | Flags | | Fragment Offset | | | | | | | | | | | | | | |
| 8 | 64 | Time To Live | | | | | | | | Protocol | | | | | | | | Header Checksum | | | | | | | | | | | | | | | |
| 12 | 96 | Source IP Address | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 16 | 128 | Destination IP Address | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 20 | 160 | Options (if IHL > 5) | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |

PacketNgin Example - IP

| Bits | 0–7 | 8–15 | 16–23 | 24–31 |
|------|----------------|------|----------|-------|
| 0 | Type | Code | Checksum | |
| 32 | Rest of Header | | | |

PacketNgin Example - IP

```
if(endian16(ether->type) == ETHER_TYPE_IPv4) {
    IP* ip = (IP*)ether->payload;

    if(ip->protocol == IP_PROTOCOL_ICMP && endian32(ip->destination) == address) {
        ICMP* icmp = (ICMP*)ip->body;
        // Echo reply
        if(icmp->type == 8) {
            icmp->type = 0;
            icmp->checksum = 0;
            icmp->checksum = endian16(checksum(icmp, packet->end - packet->start - ETHER_LEN - IP_LEN));

            ip->destination = ip->source;
            ip->source = endian32(address);
            ip->ttl = endian8(64);
            ip->checksum = 0;
            ip->checksum = endian16(checksum(ip, ip->ihl * 4));

            ether->dmac = ether->smac;
            ether->smac = endian48(ni->mac);

            ni_output(ni, packet);
            packet = NULL;
        } else if(icmp->type == 0) {
            clock_t time = clock() - *(clock_t*)(icmp->body);
            printf("Echo reply received. time : %d.%d ms\n", time/1000, time%1000);
            packet = NULL;
        }
    }
}
```


PacketNgin Example - IP

```
Starting VM...
true
***** vmid=1 thread=0 standard Output *****
Thread 0 bootting
Echo request sent
Echo reply received. time : 25.708 ms
Echo request sent
Echo reply received. time : 0.101 ms
Echo request sent
Echo reply received. time : 0.88 ms
Echo request sent
Echo reply received. time : 20.313 ms
Echo request sent
Echo reply received. time : 23.993 ms
Echo request sent
Echo reply received. time : 20.345 ms
Echo request sent
Echo reply received. time : 16.600 ms
Echo request sent
Echo reply received. time : 0.55 ms
Echo request sent
Echo reply received. time : 20.169 ms
Echo request sent
Echo reply received. time : 7.629 ms
Echo request sent
Echo reply received. time : 0.53 ms
```

PacketNgin Example - TCP

| Offsets | Octet | 0 | | | | | | | | 1 | | | | | | | | 2 | | | | | | | | 3 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---------|-------|-------------------------------------------------------------------------------------|---|---|---|-------------------|---|---|--------|---|---|----|----|----|----|----|----|-----------------------------|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|
| Octet | Bit | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 | 25 | 26 | 27 | 28 | 29 | 30 | 31 | | | | | | | | | | | | | | | | | | | | | | | |
| 0 | 0 | Source port | | | | | | | | | | | | | | | | Destination port | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 4 | 32 | Sequence number | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 8 | 64 | Acknowledgment number (if ACK set) | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 12 | 96 | Data offset | | | | Reserved 0 0 0 | | | N S | C | E | U | A | P | R | S | F | Window Size | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 16 | 128 | Checksum | | | | | | | | | | | | | | | | Urgent pointer (if URG set) | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 20 | 160 | Options (if <i>data offset</i> > 5. Padded at the end with "0" bytes if necessary.) | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| ... | ... | ... | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |

- Data offset - Header length / 4
- NS, CWR, ECE - Explicit Congestion Notification
- URG - Urgent, ACK - Acknowledgement, PSH - Push to application right now
- RST - Reset the connection, SYN - Synchronize the seq number, FIN - Close
- Window Size - Sender's buffer size (generally)
- Urgent pointer -

Source: http://en.wikipedia.org/wiki/Transmission_Control_Protocol

PacketNgin Example - TCP

```
if(ip->protocol == IP_PROTOCOL_TCP) {
    TCP* tcp = (TCP*)ip->body;

    if(endian16(tcp->destination) == 80) {
        printf("Inbound : source=%u, destination=%u, seq=%u, ack=%d, syn=%d, fin=%d\n",
            endian16(tcp->source), endian16(tcp->destination), endian32(tcp->sequence), tcp->ack, tcp->syn,
            tcp->fin);

        ether->dmac = endian48(linux_mac);
        ether->smac = endian48(ni->mac);

        ip->source = endian32(address);
        ip->destination = endian32(linux_address);
        ip->checksum = 0;
        ip->checksum = endian16(checksum(ip, ip->ihl * 4));

        tcp_source_temp = tcp->source;
        tcp->destination = endian16(8080);
        tcp->checksum = 0;
        tcp_pack(packet, endian16(ip->length) - ip->ihl * 4 - TCP_LEN);

        ni_output(ni, packet);
        packet = NULL;
    } else if(endian16(tcp->source) == 8080) {
```


PacketNgin Example - TCP

```
        onkeyup="searchBox.OnSearchFieldChange(event)"/>
    </span><span class="right">
        <a id="MSearchClose" href="javascript:searchBox.CloseResultsWindow()">
```

전 망

전 망 1 - OpenFlow의 한 계

- H/W OpenFlow Switch - 1.1 ~ 1.3
- S/W OpenFlow Switch - 1.1 ~
- OpenFlow 2.0?

전 망 2 - NFV의 한 계

- NVF Reference Implementation - OPNFV
- KVM
- Network Throughput and Response time

전 망 3 - SDN과 NFV의 결합

- OpenFlow는 Flow Control을 위한 표준 API로 (not for performance)
- NFV는 SDN의 확장 기능으로

전망 4 - 말랑말랑한 네트워크

- H/W 보다는 S/W 중심의 네트워크
- Flow Control은 OpenFlow로
- Packet Processing은 NFV(또는 유사 기술)로
- VM Migration
- Cluster Migration

Thank you for ...

Thank you!

(주)구름네트웍스
김성민 semih@gurum.cc